

Kubernetes at Datadog the *very* hard way

Laurent Bernaille and Rob Boll



Table of Contents

Background

Platform challenges

Shooting ourselves in the foot

Future plans

Background

New instance of Datadog

- Completely independent and isolated
- Launch on a different cloud provider
- Fresh start, leave legacy behind

Background

New platform requirements

- Small infra team, high leverage, low touch
- Support for multiple cloud providers
- Self service, API driven, automation friendly
- Meet our scale now, and years from now

Why Kubernetes?

Kubernetes hits all the requirements

- Extensible & API driven (but *changing fast*)
- Large active community (but *immature*)
- Scalable architecture (but *unproven at scale*)
- Multiple cloud providers supported (*kind of*)

Why Kubernetes?


Dogfooding!



+



kubernetes_cluster:app3



Hope is not an option

Platform challenges

Certificates

Runtime

Networking

Cloud integrations

Ecosystem

Scale

Certificates

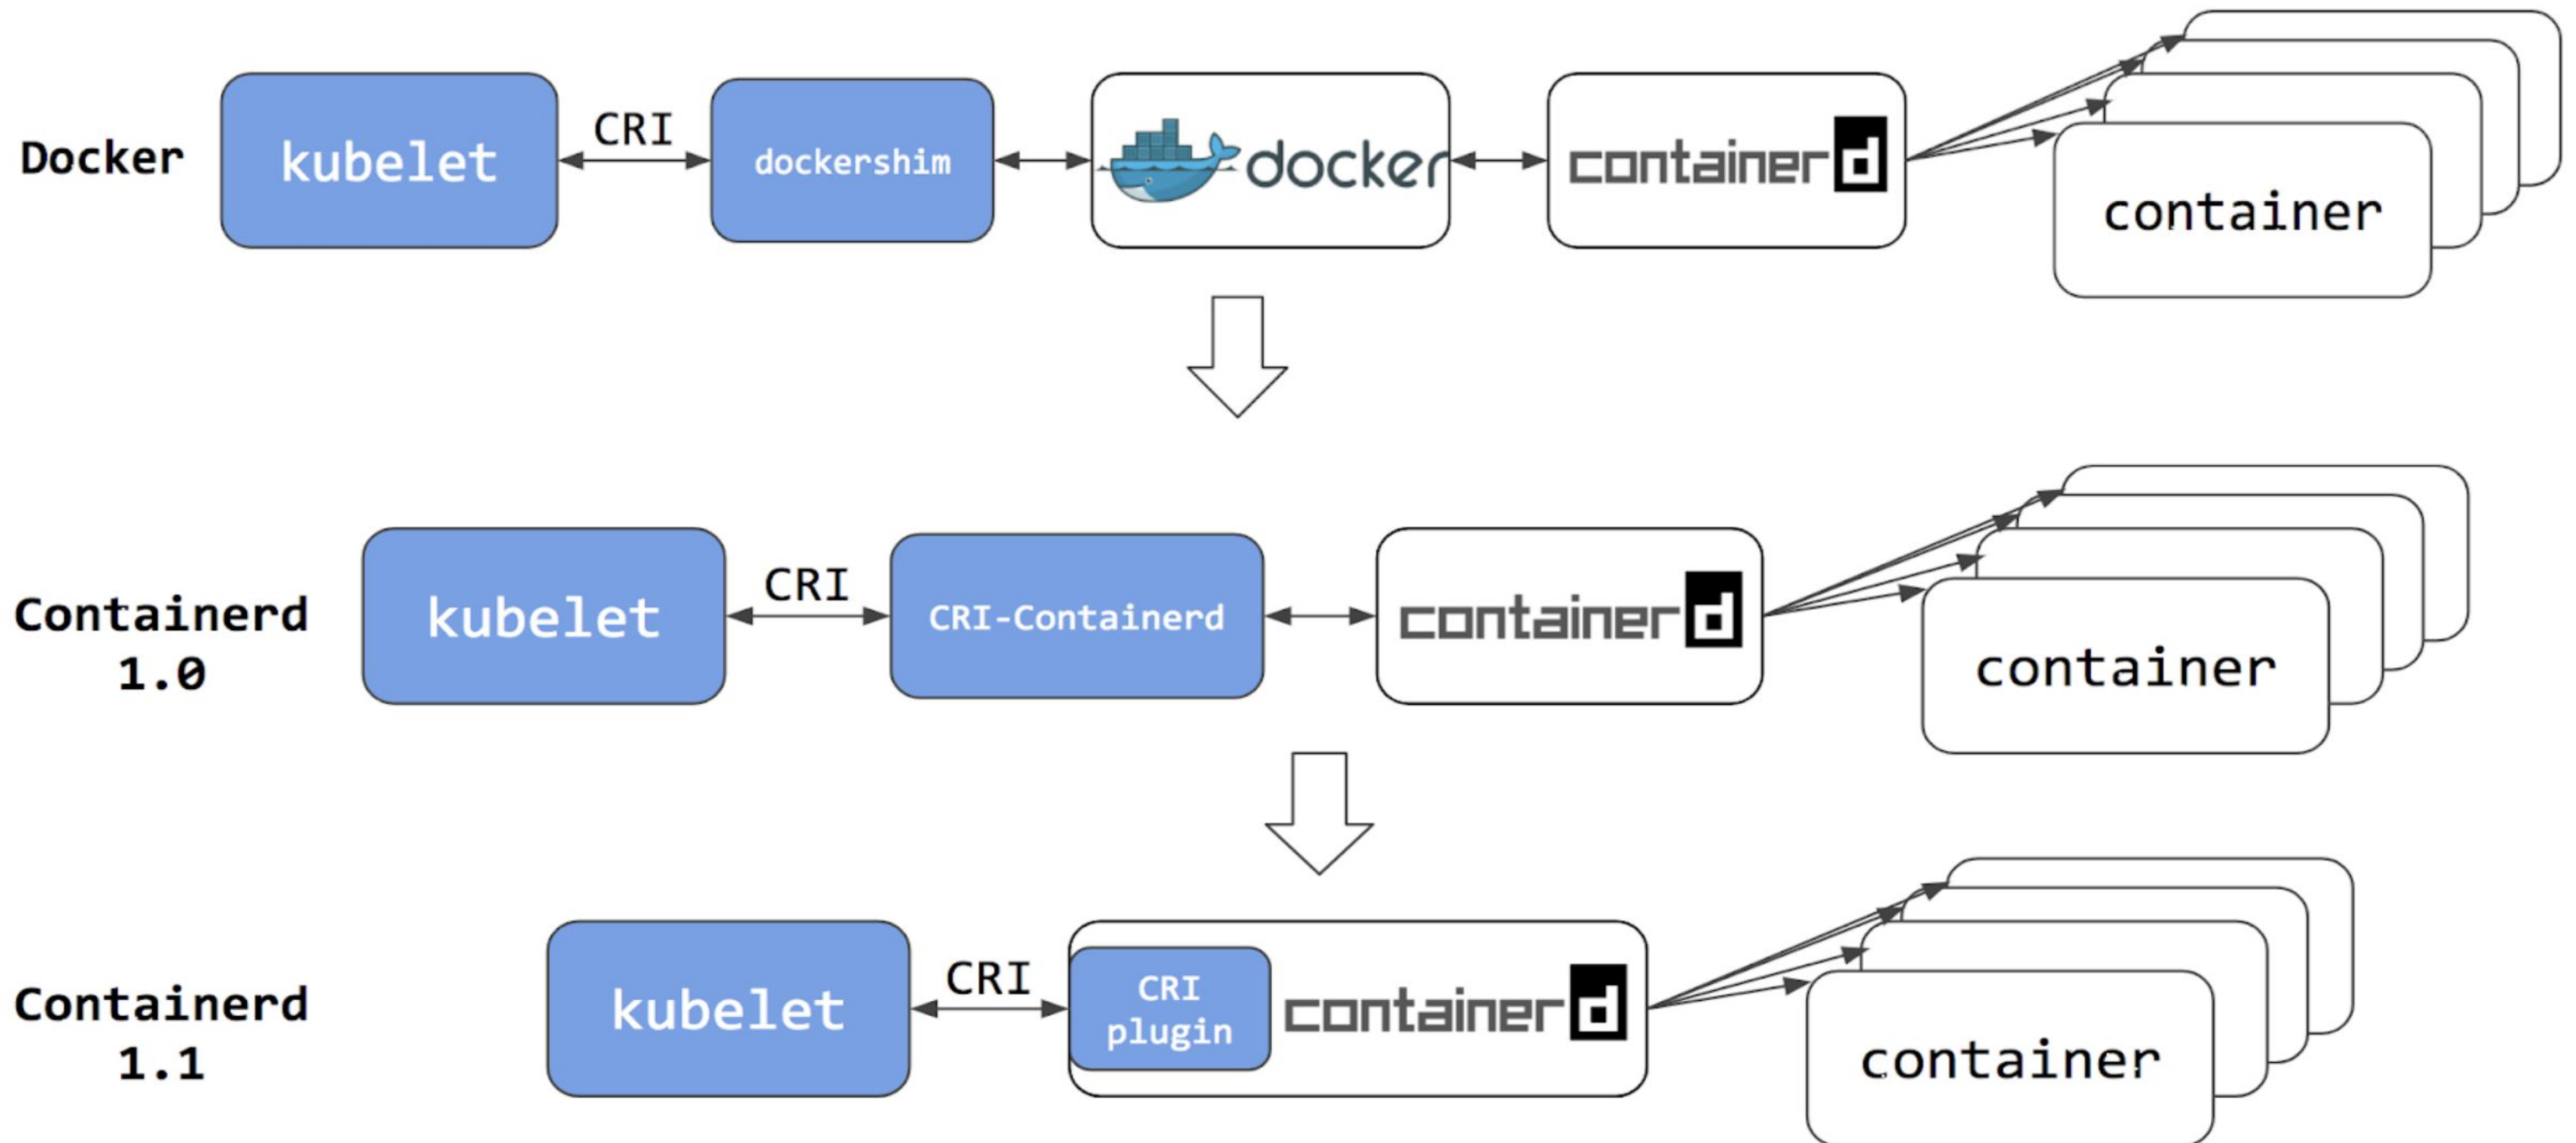
Setup

- Vault + TLS bootstrap
- Refresh certificates every 24h

The fun part

- etcd did not reload certs for client connections using IP addresses
- Kubernetes master components don't reload certificates
- Flaky bootstraps (vault dependency)
- vault-agent and vault-sidekick

Runtime: containerd



Runtime: containerd

The good

- Lightweight
- Great development team easily accessible

The bad

- Not as battle-tested as docker
- Several small issues
- Many tools assume docker (Datadog agent used to)

The ugly

- Remaining issue: shim sometimes hang and require **kill -9**

Not containerd specific...

`/home/kubernetes/bin/health-monitor.sh`

```
# We simply kill the process when there is a failure. Another systemd service will
# automatically restart the process.
function docker_monitoring {
    while [ 1 ]; do
        if ! timeout 60 docker ps > /dev/null; then
            echo "Docker daemon failed!"
            pkill docker
            # Wait for a while, as we don't want to kill it again before it is really up.
            sleep 120
        else
            sleep "${SLEEP_SECONDS}"
        fi
    done
}
```

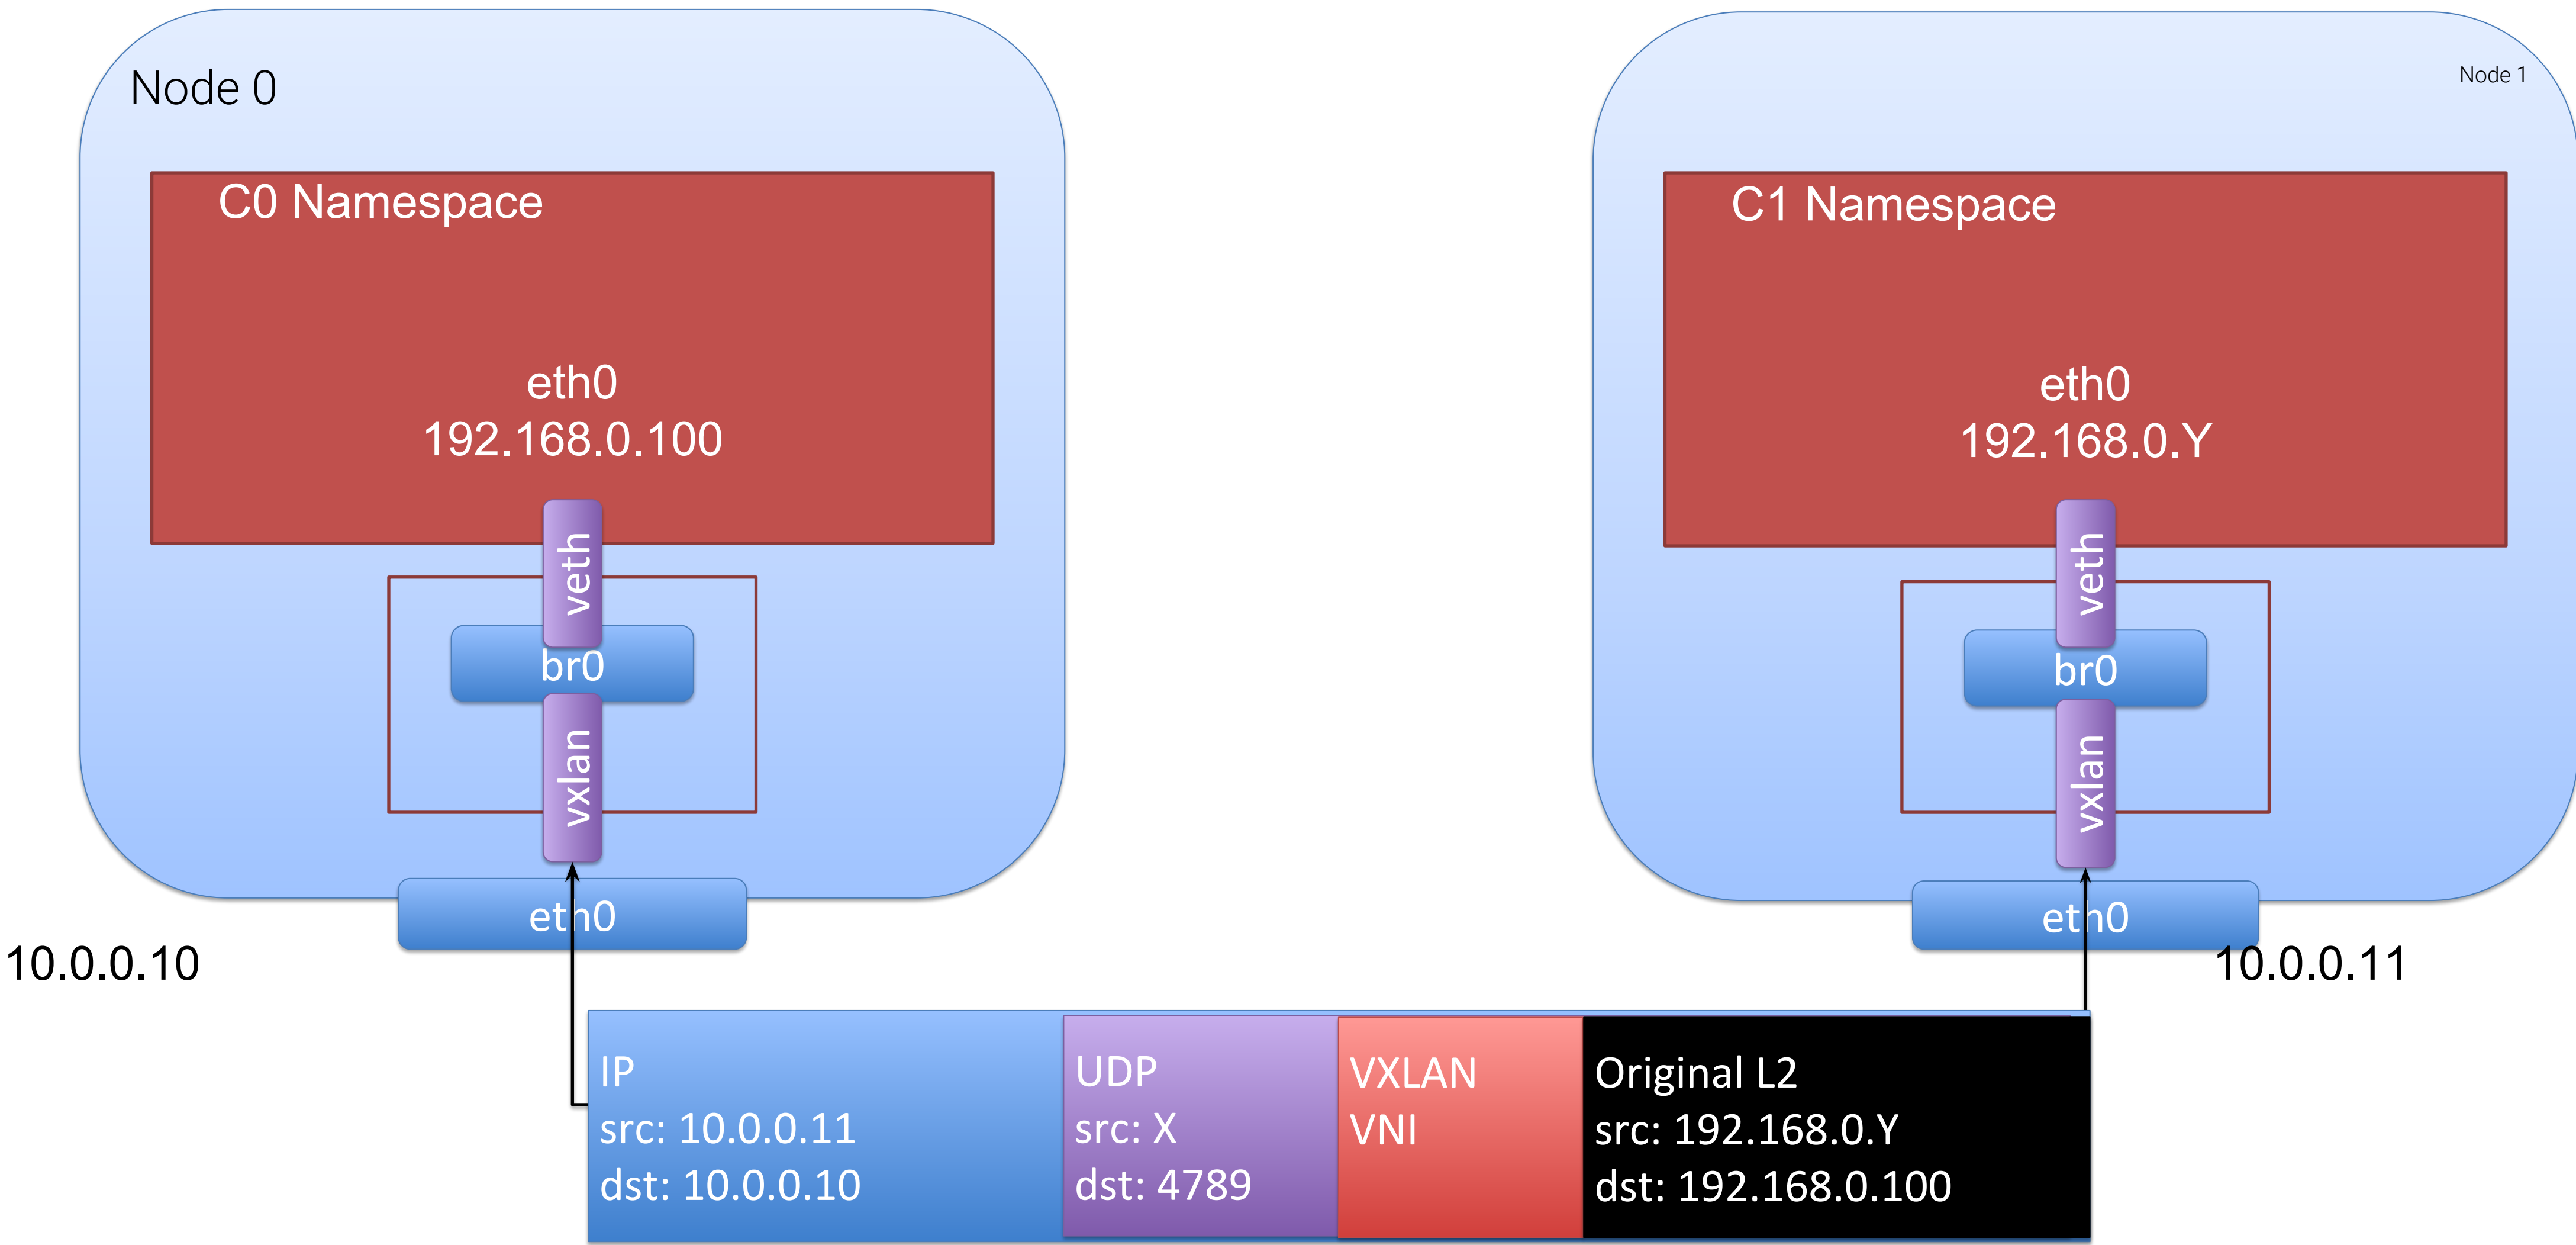
```
function kubelet_monitoring {
    echo "Wait for 2 minutes for kubelet to be functional"
    # TODO(andyzheng0831): replace it with a more reliable method if possible.
    sleep 120
    local -r max_seconds=10
    local output=""
    while [ 1 ]; do
        if ! output=$(curl -m "${max_seconds}" -f -s -S http://127.0.0.1:10255/healthz 2>&1); then
            # Print the response and/or errors.
            echo $output
            echo "Kubelet is unhealthy!"
            pkill kubelet
            # Wait for a while, as we don't want to kill it again before it is really up.
            sleep 60
        else
            sleep "${SLEEP_SECONDS}"
        fi
    done
}
```

Running on all GKE instances


<= Restart docker if docker ps hangs 60s

<= Restart kubelet if healthz takes 10+s

Networking: Overlays



Networking: Native pod routing



Networking: Native pod routing

Objectives

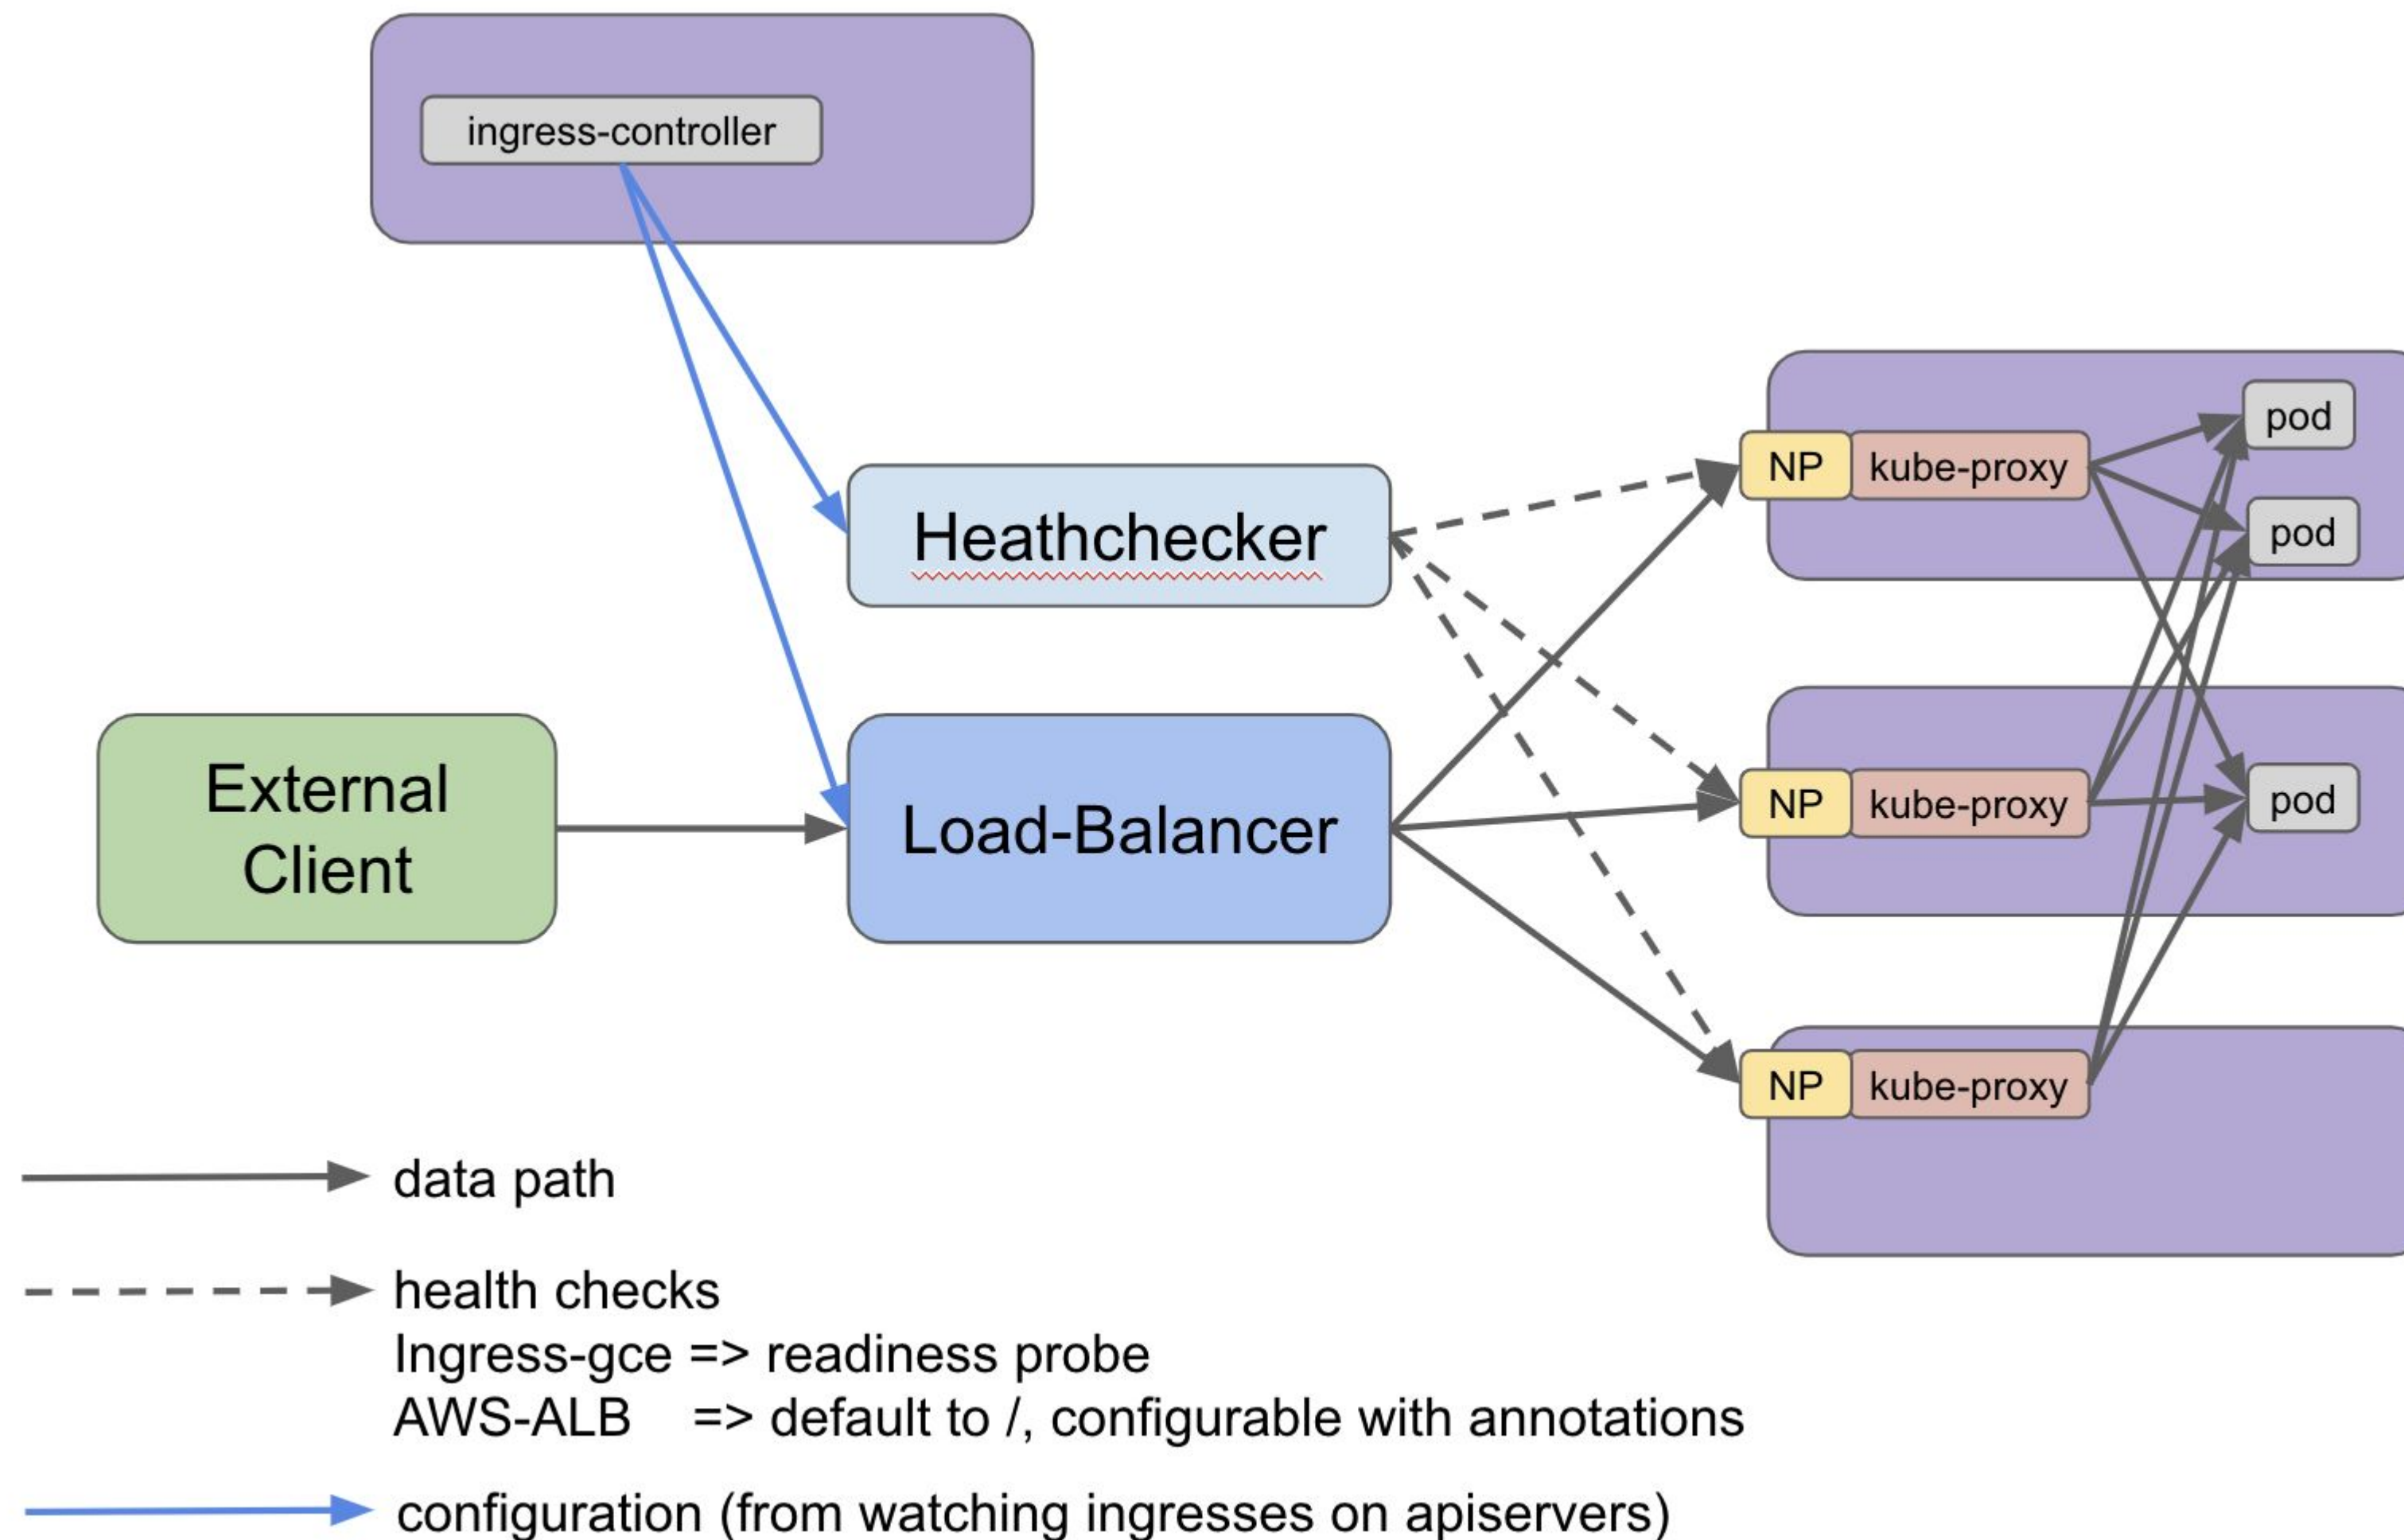
- Avoid overlays
- Avoid bridges (PTP or IPVLAN)
- Route from non kubernetes hosts/between clusters

Challenges

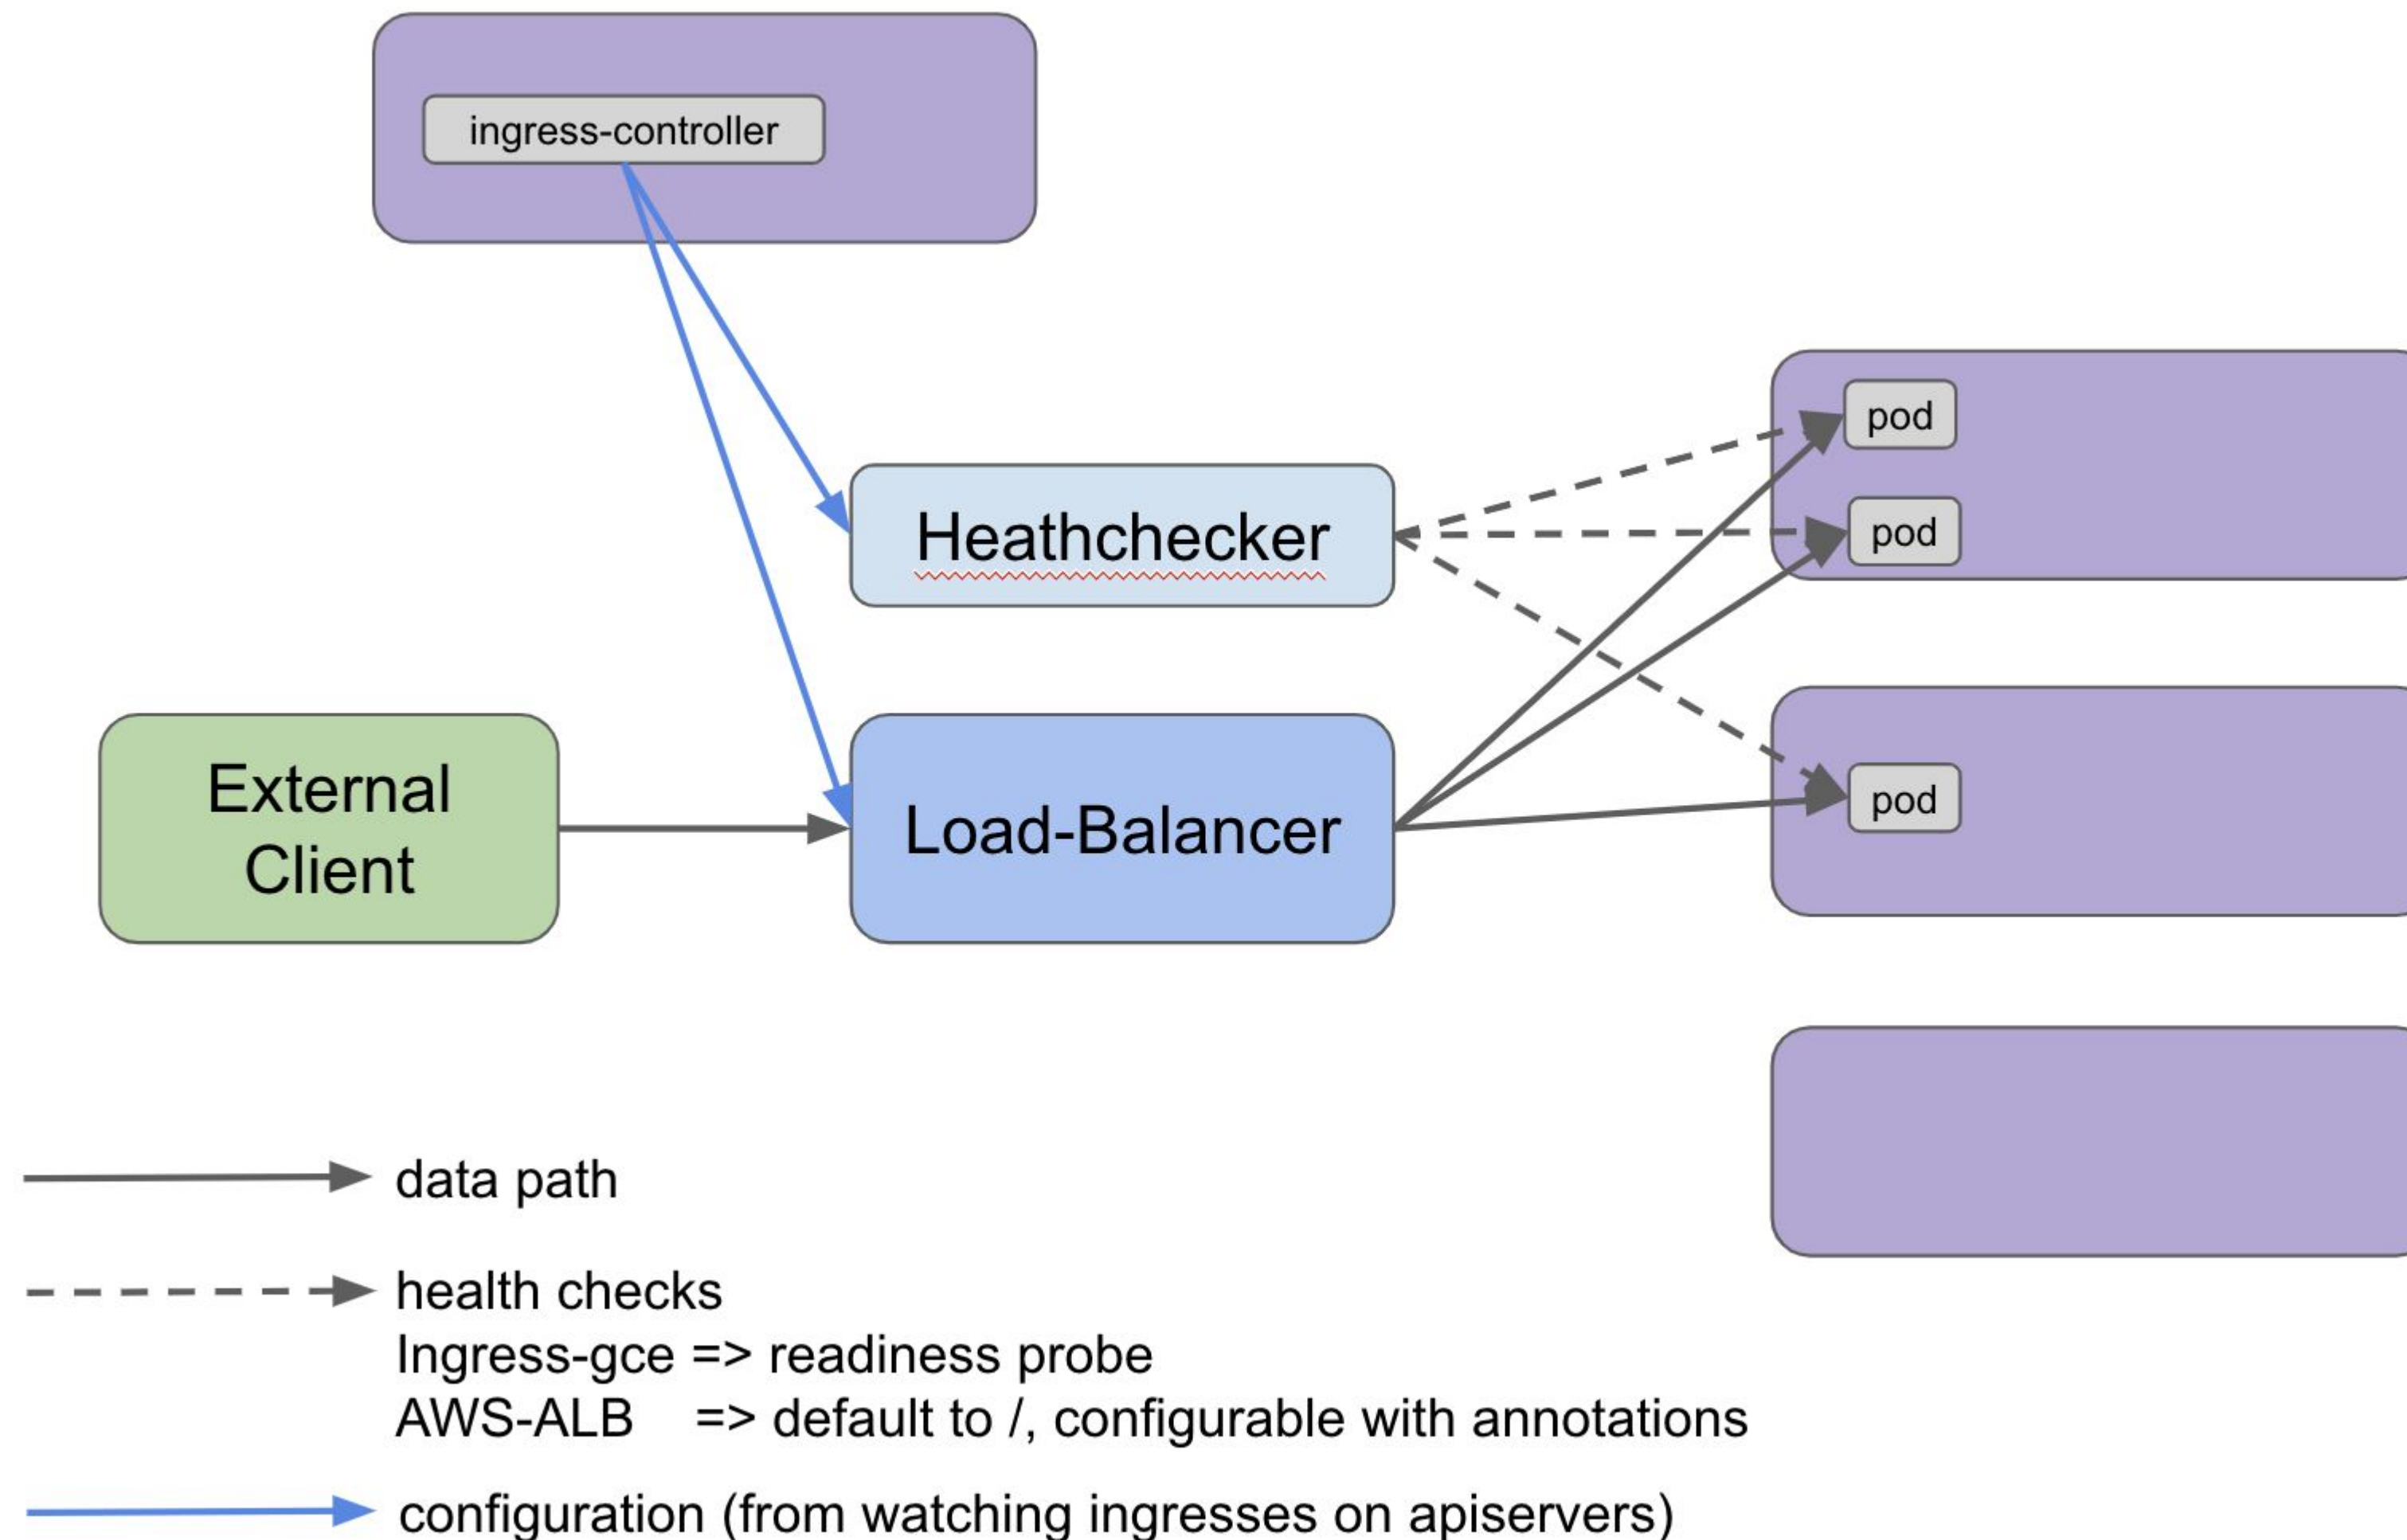
- Beta features
- Young CNI plugins
 - Bugs (Nodeports, inconsistent metadata)
 - Much more complicated to debug

➤ Good relationship with developers of the Lyft plugins

Networking: ingresses, default



Ingresses: native pod routing



Ingresses: native pod routing

More efficient


- No need to add all nodes to load-balancers
- Simpler data path

Not that simple

- Very recent
- A few bugs (instable cloud provider features, single CIDR VPC)

➤ Getting fixes upstream was quite easy

Networking: kube-proxy



IPVS instead of iptables

Sounded too good to be true

- Faster (traffic and refresh)
- Cleaner (almost no iptables rules)

Was too good to be true

- Unable to access services from the host
- Regression in 1.11 breaking ExternalTrafficPolicy: Local
- No localhost:nodeport
- No graceful termination

- Very good relationship with the Huawei team
- Almost everything has been fixed, working great so far

Networking: IPV6 and DNS

"Sometimes my DNS queries take more than 5s/time out"

"Yeah right"

[narrator]: "*Well actually...*"

- Race condition in the conntrack code
- We disabled IPV6 in the kernel
- Also had to disable native Go resolution

Cloud integrations

Many small edge-cases

- Different Load-balancer behaviors
- Magically disappearing instances (zones / instance state)
- Some “standard” controller config like “cidr-allocator-type”

[controller/nodeipam/ipam/cloud_cidr_allocator.go](#)

```
gceCloud, ok := cloud.(*gce.GCECloud)
if !ok {
    err := fmt.Errorf("cloudCIDRAllocator does not support %v provider", cloud.ProviderName())
    return nil, err
}
```

(almost) No doc

[providers/aws/aws.go](#)

```
546 //The aws provider creates an inbound rule per load balancer on the node security
547 //group. However, this can run into the AWS security group rule limit of 50 if
548 //many LoadBalancers are created.
549 //
550 //This flag disables the automatic ingress creation. It requires that the user
551 //has setup a rule that allows inbound traffic on kubelet ports from the
552 //local VPC subnet (so load balancers can access it). E.g. 10.82.0.0/16 30000-32000.
553 DisableSecurityGroupIngress bool
554
555 //AWS has a hard limit of 500 security groups. For large clusters creating a security group for each ELB
556 //can cause the max number of security groups to be reached. If this is set instead of creating a new
557 //Security group for each ELB this security group will be used instead.
558 ElbSecurityGroup string
```


Ecosystem


Good news

- Very dynamic
- Usually easy to get PRs merged


Bad news

- Not heavily tested / Limited to basic setup
- (very) Limited doc: [localVolumeProvisioner](#) and mount paths
- **Almost never tested on large clusters**
 - [cluster-autoscaler](#) doesn't work with more than [50 ASGs](#)
 - we abandoned [kubed](#) when its memory usage reached [10+GB](#)
 - [metrics-server](#): doesn't start with a single NotReady node
 - [kube-state-metrics](#): pod/node collector generates ~100MB payloads
 - [voyager](#): map sort bug lead to continuous pod recreation


Scaling nodes: 100 => 1000



API server : High load but ok
Careful with File descriptors, CPU, Memory
TargetRAM helped a lot (avoid OOM)




Controller/Scheduler with high load too
Competing for CPU => split
Thinking about splitting controllers but this is hard/impossible



etcd imbalance => shuffle etcd endpoints on API servers
works pretty well, alternatives were a bit scary (gRPC proxy)

Scaling nodes: 100 => 1000

Routes on API servers (registration + daemonsets)




CoreDNS issues

- Not enough nodes for coredns pods: "nodesPerReplica":16
- Memory limits leading to OOMkills (mem usage with “pods: verified”)


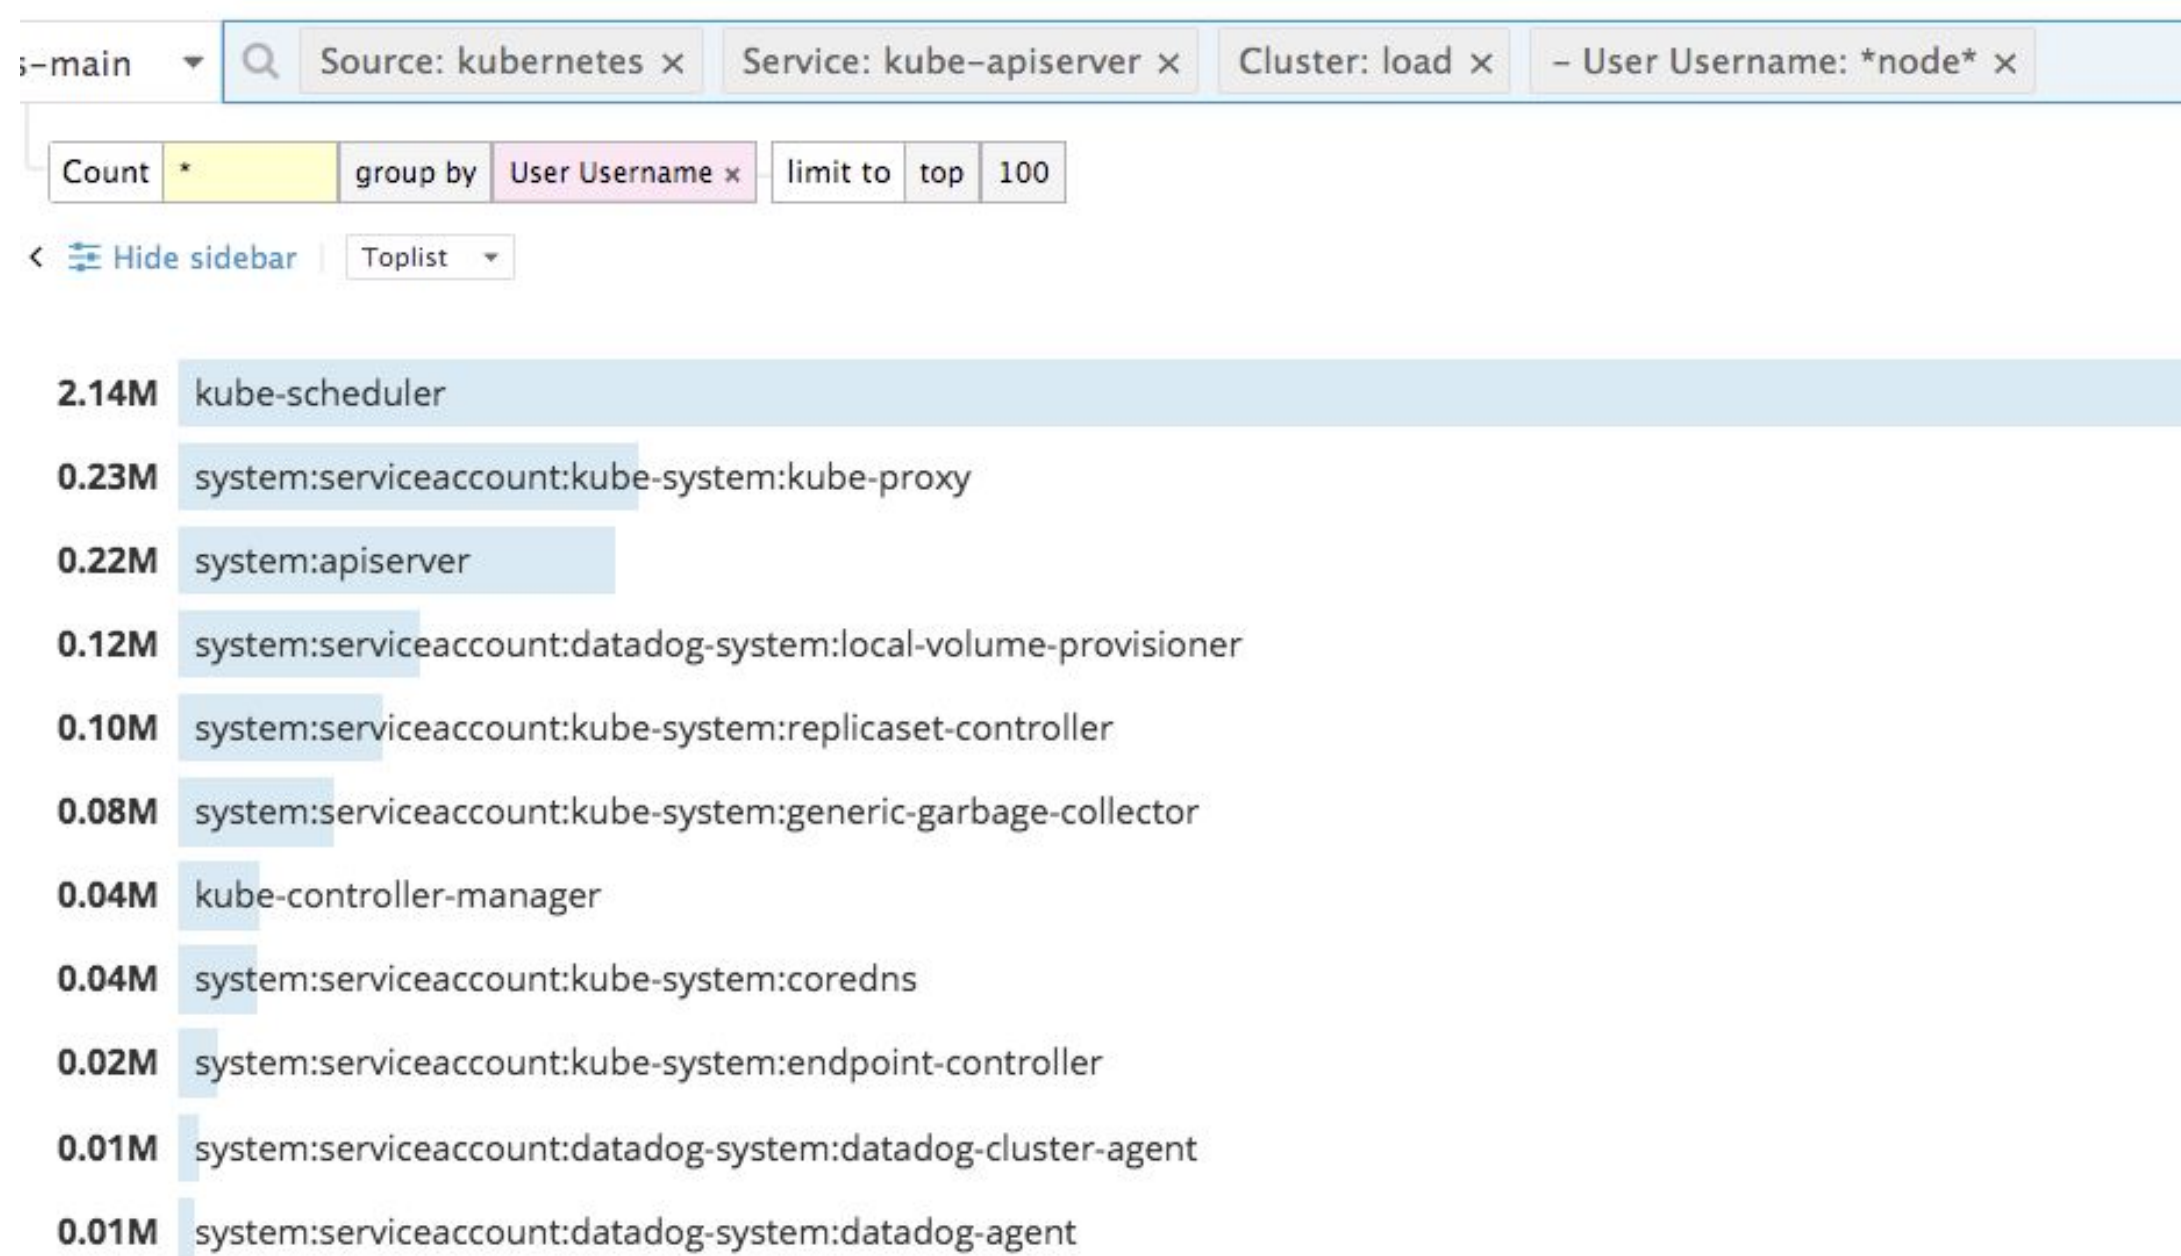
NAME	READY	STATUS	RESTARTS	AGE
coredns-7b4d675999-22d4q	0/1	CrashLoopBackOff	40	6h
coredns-7b4d675999-21t5w	0/1	ImagePullBackOff	135	17h
coredns-7b4d675999-45s94	0/1	CrashLoopBackOff	41	6h
coredns-7b4d675999-4dfbt	0/1	CrashLoopBackOff	140	17h

Scale: create 200 deployments

Very hard on all components (apiservers, controller, scheduler, etcd)



Maxed-out the scheduler QPS (too many tunables)



Main issues: apiserver sizing and traffic imbalance

Footguns

DaemonSets

StatefulSets

Cargo culting

Zombies

Containers, not VMs

Rolling updates

OOM

InitContainers

Native resources / external config

Manual changes

Daemonsets

High DDOS risk

- API servers, vault
- Cloud provider API rate-limits
- Very slow or very dangerous

Pod scheduling

- Stuck rollout
- 1.12 scheduler alpha...

Images pulled by containerd




Image pulls failed with 429 (Too many requests)



app broke due to permission change
container in restart loop
ImagePullPolicy: Always

Stateful sets

Persistent volumes

- Local: Cloud provider disk errors
- Local: New node with same name
- EBS scheduler



Laurent Bernaille @lbernail · Sep 26

Some days you know things are going to be weird^Winteresting:
cat /proc/28019/wchan
__refrigerator

localvolumeprovision: [discovery.go](https://github.com/kubernetes/kubernetes/blob/master/pkg/volume/local/local_volume_provisioner.go)

```
func generatePVName(file, node, class string) string {  
    h := fnv.New32a()  
    h.Write([]byte(file))  
    h.Write([]byte(node))  
    h.Write([]byte(class))  
    // This is the FNV-1a 32-bit hash  
    return fmt.Sprintf("local-pv-%x", h.Sum32())  
}
```


Stateful sets

Scheduling tricks

kubectl get sts myapp

NAME	DESIRED	CURRENT	AGE
myapp	5	4	5d

kubectl get pods -lapp=myapp

NAME	READY	STATUS	RESTARTS	AGE
myapp-0	1/1	Running	0	5d
myapp-1	1/1	Running	10	6m
myapp-2	1/1	CrashloopBackoff	10	6m
[?]				
myapp-4	1/1	Running	0	5d

Cargo culting

How can I keep container running on Kubernetes?

<https://stackoverflow.com/questions/31870222/how-can-i-keep-container-running-on-kubernetes>



You could use this CMD in your Dockerfile :

48

```
CMD exec /bin/bash -c "trap : TERM INT; sleep infinity & wait"
```

Zombies

```
root      8502  0.7  0.0 11032  6200 ?        Sl   16:39   0:01  \_ containerd-shim -namespace k8s.io -workdir
/var/lib/containerd/io.containerd.runtime.v1.linux/k8s.io/0eacd7463b319a9f8423f927
root      8520  0.4  0.0  46396  5768 ?        Ssl  16:39   0:00  \_ redis-server *:6379
root     10791  0.0  0.0     0     0 ?        Z    16:39   0:00  |  \_ [server_readines] <defunct>
root     11632  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [redis-cli] <defunct>
root     12222  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [server_readines] <defunct>
root     13102  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [redis-cli] <defunct>
root     14115  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [redis-cli] <defunct>
root     14500  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [redis-cli] <defunct>
root     14893  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [server_readines] <defunct>
root     15309  0.0  0.0     0     0 ?        Z    16:40   0:00  |  \_ [redis-cli] <defunct>
root     16232  0.0  0.0     0     0 ?        Z    16:41   0:00  |  \_ [server_readines] <defunct>
root     16895  0.0  0.0     0     0 ?        Z    16:41   0:00  |  \_ [redis-cli] <defunct>
root     17248  0.0  0.0     0     0 ?        Z    16:41   0:00  |  \_ [server_readines] <defunct>
root     17876  0.0  0.0     0     0 ?        Z    16:41   0:00  |  \_ [server_readines] <defunct>
root     18512  0.0  0.0     0     0 ?        Z    16:41   0:00  |  \_ [server_readines] <defunct>
root     21932  0.0  0.0     0     0 ?        Z    16:42   0:00  |  \_ [server_readines] <defunct>
root     22648  8.5  0.0  22320  5756 ?        Rs   16:42   0:00  \_ /bin/bash /usr/local/bin/server_readiness_probe.sh
```

```
ps auxf | grep -c defunct
16018
```

readinessProbe:

exec:

command: [**server_readiness_probe.sh**]

timeoutSeconds: **1**

Takeway


- Careful with exec-based probes
- Use tini as pid 1 (or shared pid namespace)

Containers, not VMs

```
Sl  11:59  0:01  \ containerd-shim -namespace k8s.io -workdir /var
Ss  11:59  0:00  |  \ /bin/bash -c -- while true; do sleep 30; d
S   12:44  0:00  |  |  \ sleep 30
Ss+ 12:00  0:00  |  \ /bin/bash
Ss  12:00  0:00  |  \ /bin/bash
Sl+ 12:01  1:37  |  |  \ python /usr/local/share/extracts/bin/b
S+  12:45  0:00  |  |  \ /bin/sh -c psql -v ON_ERROR_STOP=1
S+  12:45  0:00  |  |  |  \ /usr/bin/perl -w /usr/bin/psql
S+  12:45  0:00  |  |  |  \ /bin/bash /usr/bin/ldd /us
S+  12:45  0:00  |  |  |  \ /bin/bash /usr/bin/ldd
R+  12:45  0:00  |  |  |  \ /bin/bash /usr/bin
S+  12:45  0:00  |  |  \ /bin/sh -c psql -v ON_ERROR_STOP=1
S+  12:45  0:00  |  |  |  \ /usr/bin/perl -w /usr/bin/psql
S+  12:45  0:00  |  |  |  \ /bin/bash /usr/bin/ldd /us
S+  12:45  0:00  |  |  |  \ /bin/bash /usr/bin/ldd
R+  12:45  0:00  |  |  |  \ /bin/bash /usr/bin
```

Complex process trees and many open files is very hard on the runtime

Native resources / external config



Synchronisation

- NodePort
- External IP

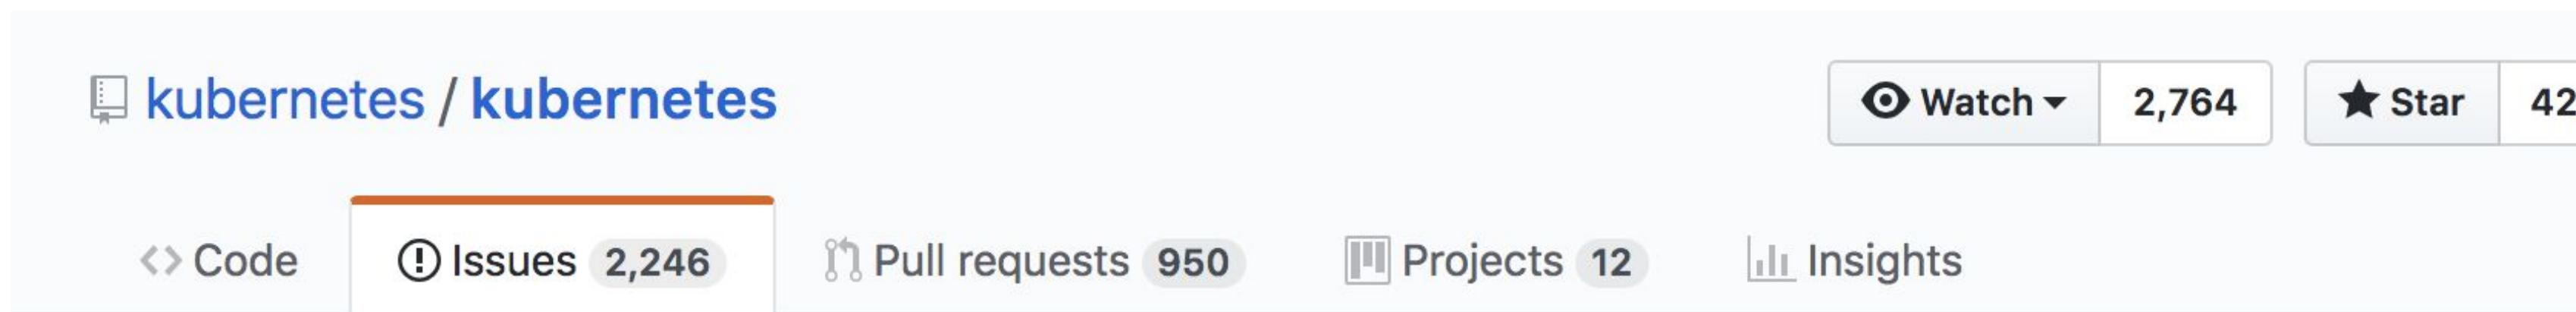
Issues

- Keeping track: Load-balancer External IP re-assigned to node
- Port conflict on Internal Load-Balancer (port 443, **broke apiservers...**)

Rolling updates

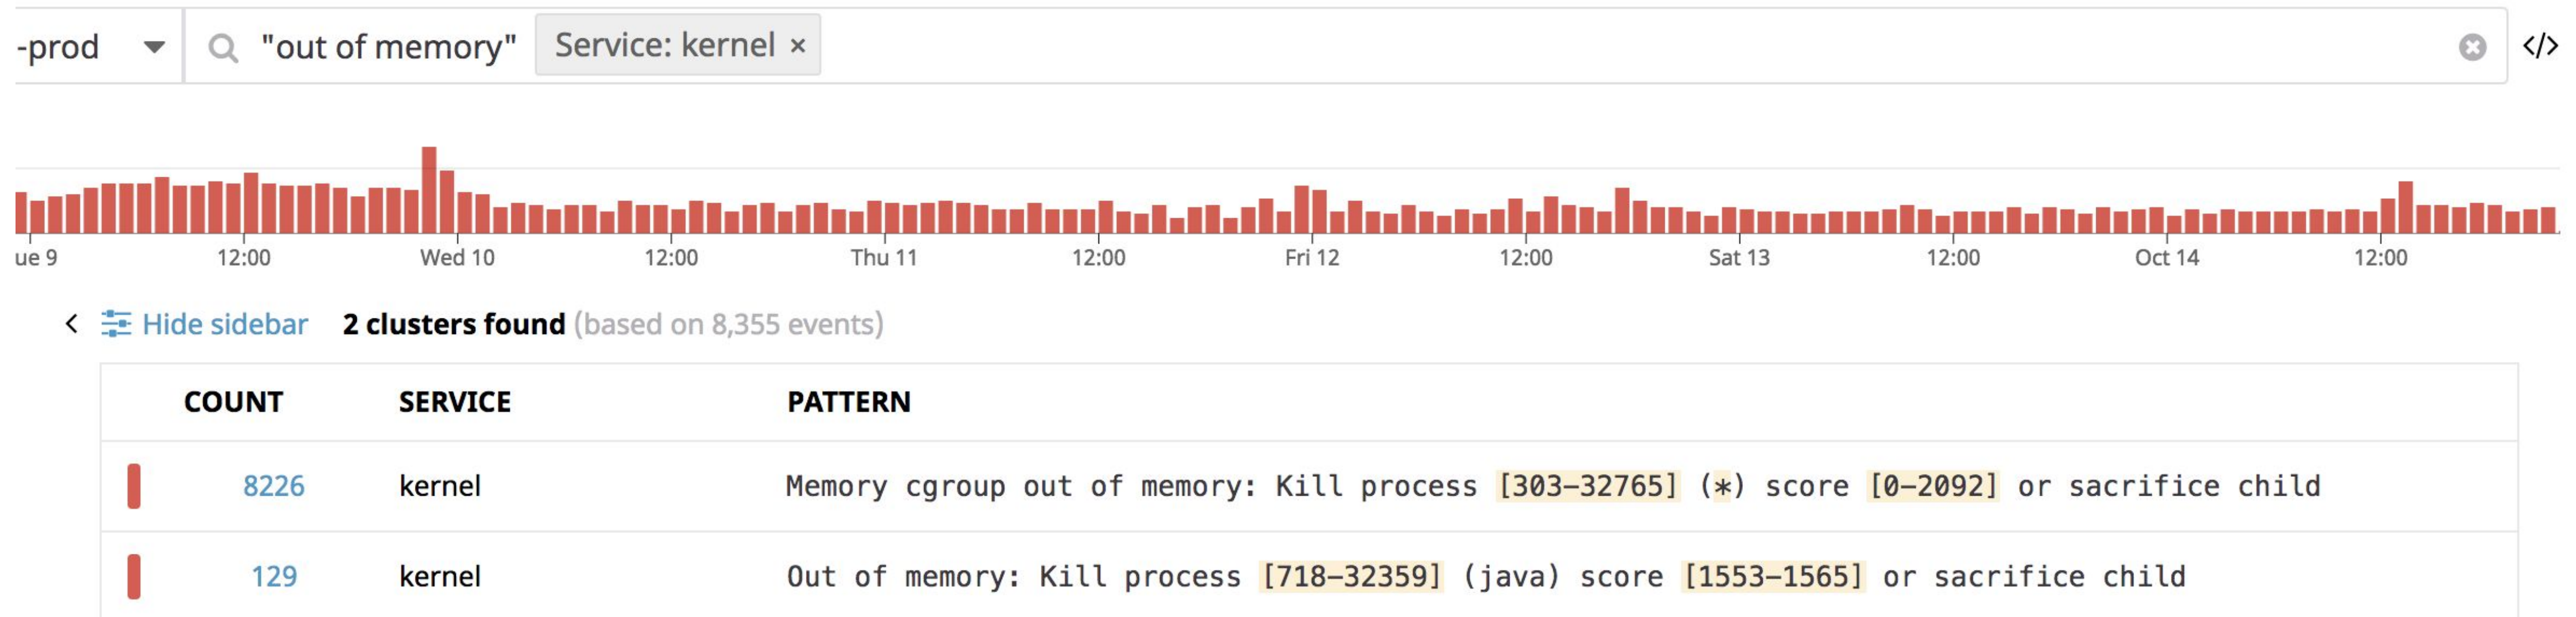
spec.replicas and hpa

- Replicas override hpa settings
- Removing replicas is not enough
- Recommended solution: edit the “last-applied” deployment first



Removing spec.replicas of the Deployment resets replicas count to single replica #67135

OOM Killer



Limits too low will trigger "cgroup oom"

Requests too low (or 0) will trigger "system oom"

Not a surprise, but do spend some effort on sizing

InitContainers

Requests/Limits

- Pod Resources = $\text{MAX}(\text{MAX}(\text{initContainers}), \text{sum}(\text{containers}))$
- LimitRanger also applies to InitContainers

Inconsistent behavior on container restarts

- *Usually* not restarted
- *Except* when Kubernetes no longer knows their exit status

Manual changes

Untracked kubectl changes

```
kubectl apply / edit
```

Partial chart apply

```
kubectl apply deploy
```

```
=> checksum/config_templates: 840ae1e0b4b2f7b5033edd34fd9eb88b55dc914adca5c
```

^ Hash for the updated config map, but not deployed

Chart deletion

```
kubectl delete -f <dir>
```

Contained namespace

Future Plans

Control plane isolation

- "Meta" cluster for control planes

Better load balancing

- Avoid imbalanced API servers
- Avoid imbalanced etcd

Container-Optimized images

- In-place upgrade for data stores

Custom controllers

- We already have a few but will build new ones

Conclusion

The Kubernetes Control plane is very complex

- Crazy number of options/tunables
- Low-level components are great but some bugs remain
- The ecosystem is very young
- Reading the code is not optional (and fixing part of it)

Account for culture change and training

- Kubernetes resources don't look that complicated
- Many, many edge-cases and pitfalls

Instrument the audit logs

- Very high value to debug performance issues
- Also helps understand history of interactions with a specific resource
- Expensive (very verbose: we are at 1000+ logs/s) but game changing for us

Thank you

Also, We're hiring

