



Monitoring Kubernetes with Prometheus

Tobias Schmidt - ContainerDays NYC November 4, 2016

github.com/grobie - [@dagrobie](https://twitter.com/dagrobie)

github.com/grobie/prometheus-on-kubernetes

Bootkube or Minikube

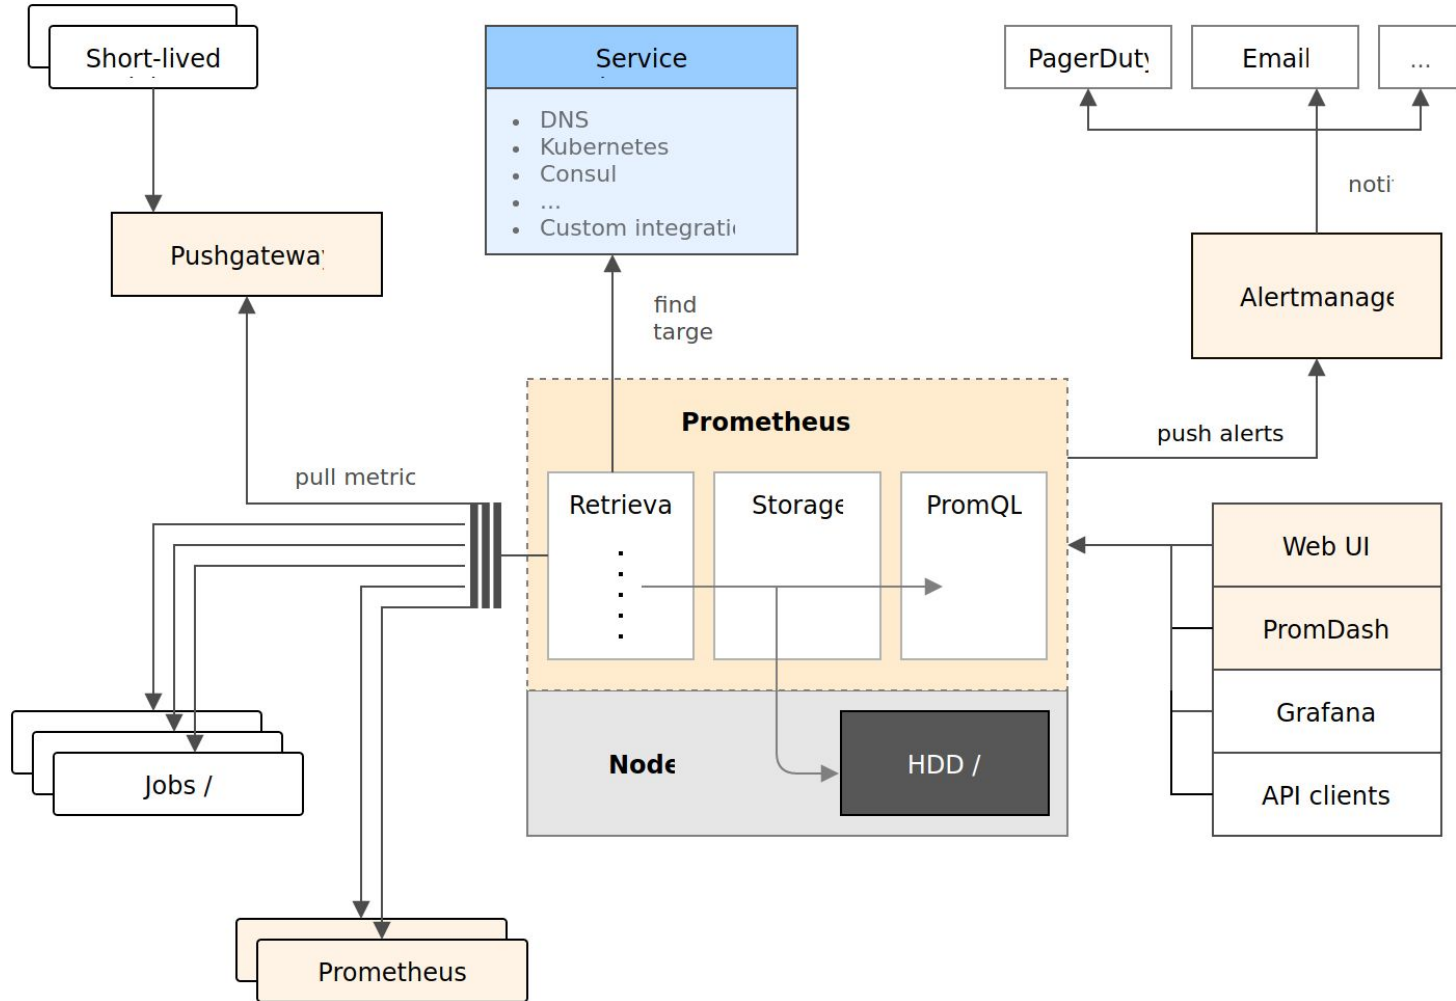
Prometheus

Monitoring system and time series database

Motivation

- Microservice architecture
 - Many more services than in traditional host-based monitoring
 - Short lifecycles
 - Heterogeneous workloads
- Insight
 - Detailed (instance/endpoint/version/... drilldown) and aggregated (across a service) of everything (hardware to service)
 - Trends (act before something becomes a problem)
- Alerting
 - Symptom vs. Cause
 - Grouping, flexible silencing

Overview



Examples

```
# HELP etcd_store_writes_total Total number of writes seen ...
# TYPE etcd_store_writes_total counter
etcd_store_writes_total{action="compareAndDelete"} 2
etcd_store_writes_total{action="compareAndSwap"} 4016
etcd_store_writes_total{action="create"} 218
etcd_store_writes_total{action="set"} 5
```

```
count by(job)(up == 0) / count by(job)(up)
rate(etcd_store_writes_total{action="set"}[1m]))
sum without(action)(rate(etcd_store_writes_total[1m]))
```

Configuration

```
# prometheus.yaml prometheus.io/docs/operating/configuration/
```

```
global:
```

```
    # Settings applying to all jobs
```

```
scrape_configs:
```

```
    # Define different scrape jobs
```

```
Rules_files:
```

```
    # Load files specifying rules to pre-calculate expressions
```

```
    # as well as alerts.
```

Configuration

scrape_configs:

- job_name: etcd

static_configs:

- targets: ["172.17.4.51:2379"]

- job_name: kube-components

kubernetes_sd_configs:

- role: endpoints

relabel_configs:

- # Custom filtering and label mapping

Configuration

```
# continued
```

```
relabel_configs:
```

- action: keep
source_labels: [__meta_kubernetes_service_name]
regex: "kube-(.*)-prometheus-discovery"
- action: keep
source_labels: [__meta_kubernetes_endpoint_port_name]
regex: "prometheus"
- action: replace
source_labels: [__meta_kubernetes_service_name]
target_label: job
regex: "(kube-.*)-prometheus-discovery"

Kubernetes

Container orchestration system

Domain objects

- Pod
 - Group of one or more containers, share context and namespaces
 - Co-located and co-scheduled (allows for side-cars)
- Service
 - Logical set of Pods, stable access points
- Deployment
 - Declaration of the desired state (what to run, how to get there)
- Daemon / Pet / Replica sets
 - Definition of groups of pods (each node / stateful / stateless)
- ConfigMap
 - Configuration data / files (can be mounted in containers)

Workshop

Monitoring Kubernetes with Prometheus

git checkout 1.setup

kubectl get nodes

Running Prometheus in Kubernetes

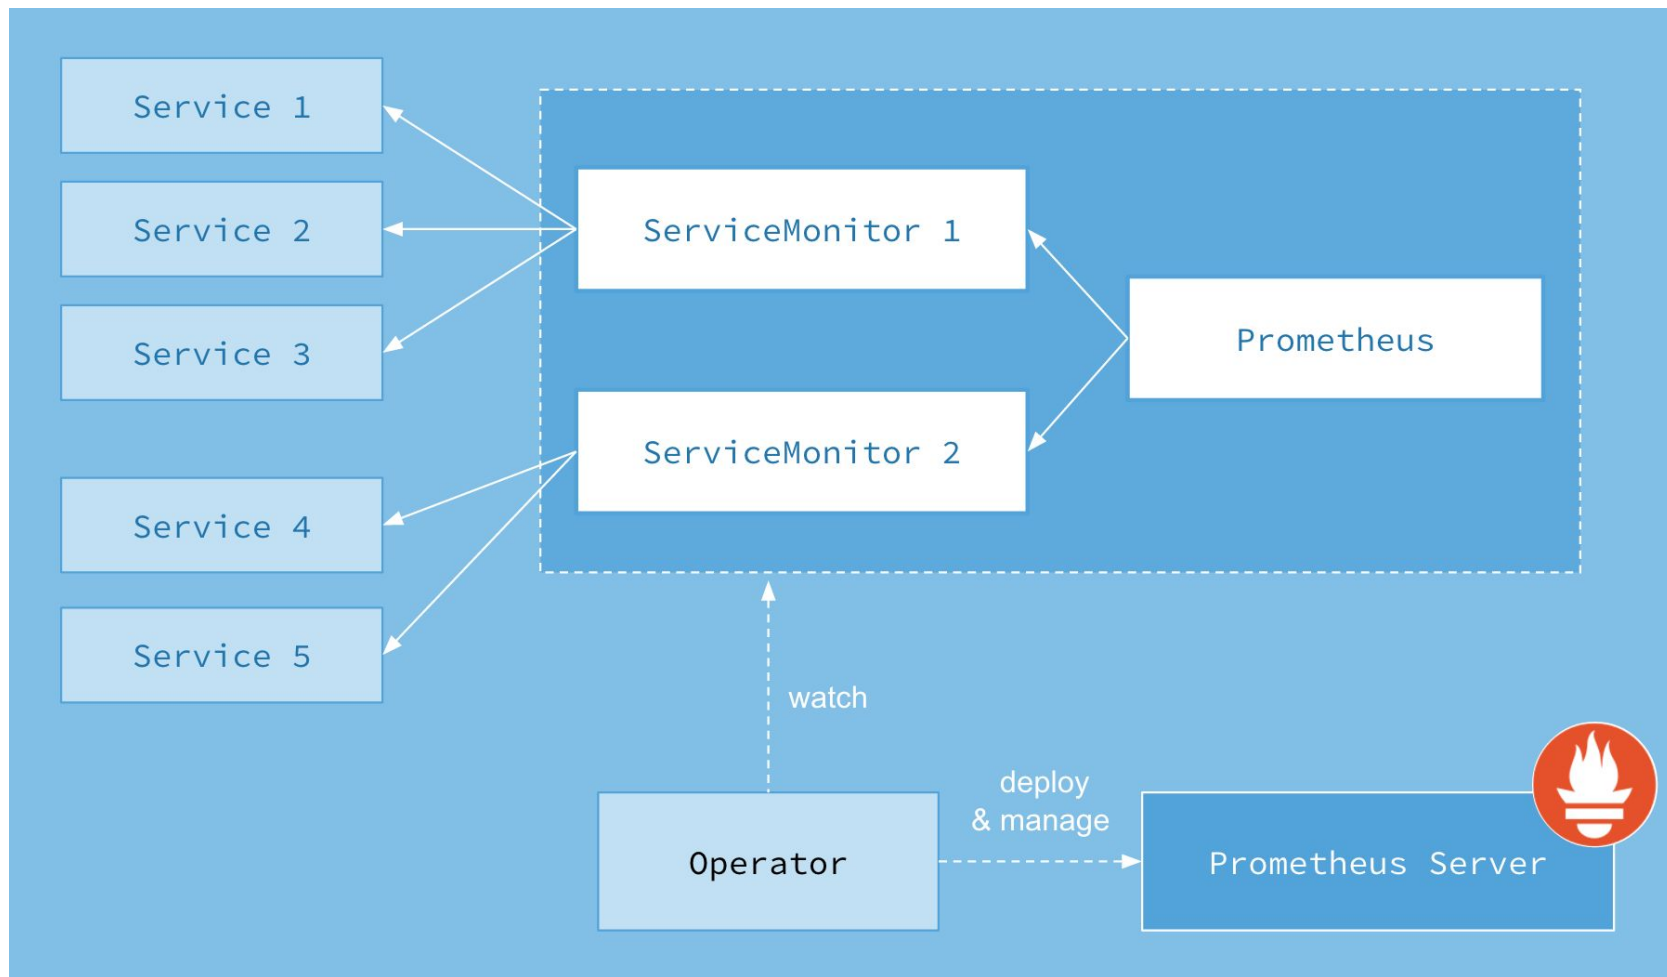
Installation and configuration

Running Prometheus

inside of Kubernetes

- What we need
 - Pod specification defining how to run Prometheus
 - Load and manage configuration
 - Service specification to access Prometheus on stable IP
- Options
 - Write own pod+service+petset+... manifests
 - Kubernetes Helm chart in the making
<https://github.com/kubernetes/charts/pull/151>
 - CoreOS wrote an Operator managing Prometheus and its configuration: <https://github.com/coreos/kube-prometheus>
<https://coreos.com/blog/the-prometheus-operator.html>

Prometheus Operator




```
git checkout 2.install-prometheus  
scripts/deploy
```

Monitoring Kubernetes infrastructure

Configuration and discovery

```
git checkout 3.monitor-nodes  
scripts/deploy
```

```
git checkout 4.monitor-kubernetes  
scripts/deploy
```

```
git checkout 5.install-grafana  
scripts/deploy
```

Monitoring services in Kubernetes

Configuration and discovery

```
git checkout 6.monitor-example-app  
scripts/deploy
```

Practical examples

Queries and dashboards


```
git checkout 7.add-rules  
scripts/deploy
```



Incoming Request Rate (HTTP & Thrift)

19K rps


99th Percentile Latency

162.6 ms

Error Rate

0.24 rps


Incoming HTTP Request Rate



Incoming HTTP Request Latency




Incoming HTTP Request 5xx Breakdown



Outgoing HTTP Request Rate




Outgoing HTTP Request Latency



Outgoing HTTP Request 5xx Breakdown




Client Failure Accruals



Client Failure Percentage




Client Retries / Requeues




Thrift Successes




Thrift Failures



Thrift 99th Percentile Latency



Further reading



My Philosophy on Alerting

Rob Ewaschuk

https://docs.google.com/document/u/1/d/199PqyG3UusyXlwi_eHaqbGiWVa8eMWi8zzAn0YfcApr8Q/preview

Thank you

Tobias Schmidt - ContainerDays NYC November 4, 2016

github.com/grobie - [@dagrobie](https://twitter.com/dagrobie)